Who Speaks to Whom? Sponsorship, Spatial Discourse, and Political Identity in French Electoral Manifestos (1973–1993)

Zakaria Bouliaire
ENSAE Paris
zakaria.bouliaire@ensae.fr

Abstract

We apply Named Entity Recognition to text from the Archelec corpus, a dataset consisting of digitised French legislative electoral manifestos at various elections between 1973 and 1993. We pursue two questions. First, which national politicians did local candidates call to mind, and does that selection suggest alignment or distancing? Second, do candidates orient their rhetoric locally or toward national and international stakes? Extracting person names and geographic locations from the 2,551 candidate manifestos using spaCy's `fr_core_news_lg` model. OCR-induced variants are resolved through a canonical name normalisation pipeline and the tone of each citation is approximated using a lexicon-based context-window proxy. The spatial analysis classifies spaCy-detected geographic entities as local, national or international using a curated reference list. We show that Mitterrand completely dominates citation counts for all parties, peaking in 1981 and 1988 following his presidential victories, and that national references consistently overwhelm local ones over time. It turns out that the two rhetorical dimensions appear to be largely independent of one another. These results should be interpreted as exploratory, given OCR noise and the limited geographic coverage of the corpus.

1 Introduction

Electoral manifestos, is an excellent resource for studying political discourse. In these manifestos, we can analyse rhetorical strategies reaching well beyond their programmatic content: candidates talking about political figures, defining their positions and the ones relative to themselves, and frame policy issues locally, nationally or globally. These choices indicate the political position a candidate wishes to claim.

The Archelec corpus gives a long-term view on such strategies. In this corpus there are the professions de foi of five French legislative elections: 1973, 1978, 1981, 1988 and 1993. It straddles the end of Gaullist hegemony, the rise of the Socialist Party, 2 periods of presidential cohabitation and the emergence of the Front National. The corpus, transcribed using OCR, contains a lot of noise (e.g. recognition errors, inconsistent capitalisation and layout artefacts) making it a testbed for NLP methods on historical documents.

We address two specific questions drawn from the Archelec project call:

1. **Sponsorship**: which politicians are cited by local candidates, and to what extent do they agree with them or reject them ?
2. **Spatial discourse**: do candidates focus their manifestos in local level politic or national and global ones?

Preprint.

In both instances, the point is not to account for political stances in full detail, but rather to classify systematic designs of rhetoric within a historical corpus permeated with noise.

2 Related Work

The digital humanities have increasingly focused on named entity recognition from historical and noisy documents. An extensive survey of NER for historical newspapers can be found in Ehrmann et al. [2023] where they mention that OCR noise systematically reduces the performance of approaches trained on clean text, therefore domain adaptation is often required. This is supported by the high recall but low precision of `fr_core_news_lg` on administrative OCR documents, which comprises a benchmark dataset for French administrative documents and can be found in the AdminSet corpus Vidal et al. [2024].Since OCR transcribes electoral texts with similar noise characteristics to common vocabulary, these results are directly applicable to our setting.

In terms of modelling, the models have shifted from rule-based systems to neural sequence labelling ones. In this work, we use spaCy's statistical pipeline, which relies on transition-based parsing and word-vector features Honnibal et al. [2020]. More recent works are based on fine-tuned transformers : CamemBERT Martin et al. [2020], a French BERT model trained on the OSCAR corpus, and CamemBERT-bio have been shown to outperform baseline models on French NER benchmarks. In a promising approach for domain specific extraction tasks such as ours, zero-shot models like GLiNER Zaratiana et al. [2023] enable conditioning on entity type descriptions and eliminate the need of any labelled data.

In political science work on the Comparative Manifesto Project Budge and Klingemann [2001] demonstrated that party manifestos can be systematically coded for issue salience, but these approaches are dependent upon human annotation and so do not scale to large corpora. Somewhat later, computational methods such as topic models Quinn et al. [2010] and word embeddings Rheault and Cochrane [2020] have been used to analyze political text at scale. Approaches to political alignment based on sentiment have been used for parliament debates Abercrombie and Batista-Navarro [2020] but are still unusual in the literature of electoral manifestos.

The sponsorship question we are addressing here relates to a basic concept in political communication theory: the notion of valence (the idea that by mentioning someone who matters and is salient, one can either borrow some legitimacy from him or alternatively define himself against an opponent). For example, lexicon-based sentiment scoring approaches using the context window has been applied to study political references in news corpora Abercrombie and Batista-Navarro [2020], which we adapt here for use with electoral manifestos.

In the light of the fact that there is no labelled training data available for this corpus, our aim in this paper is not to achieve highest possible scores on NER but rather build an interpretable baseline and evaluate what we can learn from it already. This is in the pattern of first applying a pretrained NLP pipeline on some domain specific corpus followed with more complex supervised or transformer based techniques.

3 Data

3.1 The Archelec corpus

The data set is made up of text files separated by election year that were extracted from scanned profession de foi documents. The unit of analysis is therefore the candidate manifesto, not the individual page. We keep only the profession de foi documents and do away with bulletins–de–vote. This results in 2,551 documents being treated as candidate manifestos across five different elections (Table 1) after parsing and grouping.

The corpus spans five departments: Bas-Rhin, Haut-Rhin, Rhone (Rhone), Pyrenees-Orientales and Hautes-Pyrenees. The geographic scope is limited and should be considered when interpreting spatial findings. This geographic restriction means that the corpus should not be interpreted as representative of all French legislative campaigns. Candidate counts and manifesto length distributions are displayed in Figure 1.

Table 1: Corpus statistics by election year.

Year	Candidates	Mean words
1973	473	5,821
1978	478	5,934
1981	487	5,612
1988	562	6,043
1993	551	5,967
Total	2,551	5,876

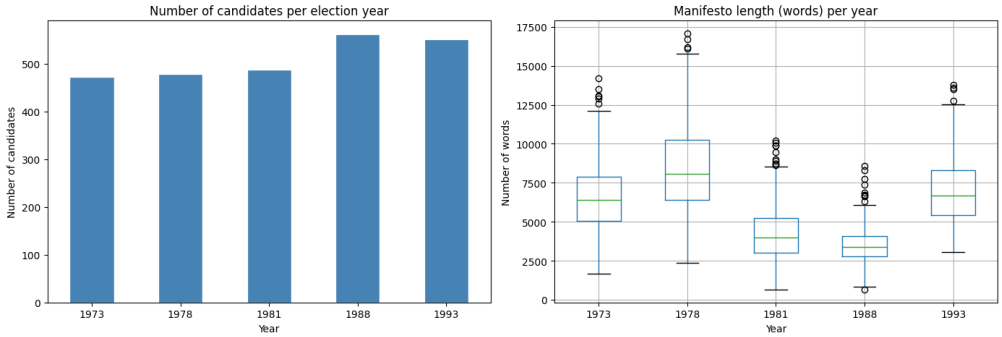

Figure 1: Left: Number of Candidates per election year. Right: distribution of manifesto length (ie number of words) by year. The corpus grows a little over time but the lengths og the manifestos fluctuate around a mean of 5,876 words.

3.2 Metadata and party classification

Candidate-level traits were downloaded in CSV format from the Archelec explorer. There are more than 100 unique values for the raw party field (`titulaire-soutien`). For the smaller parties we performed manual matching but for these nine political families mapping was mainly done using keyword matching. Party-specific analyses exclude candidates without a party mention, leaving 1,514 out of 2,551 candidates for party-level comparisons. These results should therefore be interpreted as applying to the subset of candidates for whom a party family could be assigned.

4 Methods

4.1 Named Entity Recognition

For this case study, we extract entities of type `PER` and `LOC` from each manifesto using spaCy's model fr_core_news_lg. The goal is not to benchmark NER performance exhaustively, but to use a transparent pretrained baseline and then apply task-specific post-processing to make the extracted entities usable for aggregate analysis. The raw output is noisy: an average of 84 detected person entities per manifesto, inflated due to OCR artefacts (all-caps words misclassified as named entity). We filter once to restrict us only to those entities containing at least one lowercase letter (which helps reduce OCR-related false positives), and again removing an ad-hoc stoplist of title & function names. This filter is deliberately simple: it removes many OCR-induced artefacts, but it may also discard some legitimate all-caps mentions. For case normalisation, we are using a Title Case function that takes care to handle hyphens and apostrophes correctly.

4.2 Name normalisation

OCR error & and the typographic variations that there are in the corpus(e.g. "Francois Mitterrand", "F. Mitterrand") generates the same figure in several surface forms. We construct a canonical name mapping for the top dozen most relevant individuals, and apply it across all contributors so that we can reduce fragmentation, enabling trusted aggregation of the citation data in the Archelec corpus.

4.3 Context window analysis

Every time we see a mention of a target figure, we extract a 10 word window around the mention and keep positive features (soutien, fidèle, favorable, accord and victoire) against negative ones (contre, opposition,danger,échec and dénonçons). The alignment score is:

$$s = \frac{n_{pos} - n_{neg}}{n_{pos} + n_{neg}} \in [-1, +1]$$

A score of $+1$ indicates pure endorsement; -1 indicates pure distancing. Only citations with at least one marker in context are included. This score should be interpreted as an exploratory lexical proxy for alignment or distancing, rather than as a definitive measure of political endorsement.

4.4 Spatial classification

We distinguish three types of geographic references: local (in our 5 corpus departments and their main cities), national (France, Paris, the Republic) or international (Europe, abroad). We classify spaCy-detected LOC entities using a curated list of local, national and international references. This approach remains dependent on the quality of spaCy's LOC extraction, which may also suffer from false positives in this corpus.

5 Results

5.1 Who gets cited, and when?

Mitterrand dominates citation counts, far ahead of any other figure (this dominance should not be interpreted as endorsement: high citation counts may reflect either support, opposition, or the centrality of a figure in national political debate), with 6149 total mentions next after Laguiller (1480), De Gaulle (1403) and Rocard (1278). Citation frequency from Figure 2 and normalised by the number of candidates for each year. Mitterrand spikes sharply in 1981 and again in 1988, after his presidential victories, and then falls sharply in 1993. General de Gaulle, who died in 1970 was still extensively cited as late as 1973 and 1978. Le Pen is shown progressively (from 1988) with the rise of Front National.

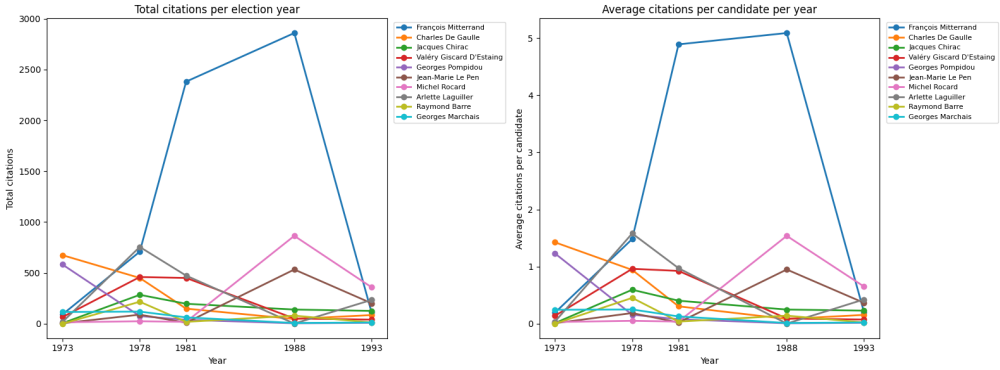

Figure 2: Average number of citations per each candidate per election year for the 10 most cited national figures. Mitterrand peaks in 1981 and 1988 in line with his presidential victories. De Gaulle declines steadily even though he died in 1970. Le Pen emerges from 1988.

5.2 Who cites whom?

The average number of citations per candidate by political party is shown in figure 3. Mitterrand is cited across the political spectrum, from RPR/Gaulliste candidates to FN candidates, who mention him 2.4 times per candidate on average, most likely as part of opposition rhetoric. Yet portrayals continue to cling on De Gaulle as a reference even for parties including PS/Gauche, since he is still

remembered as one of the founders and icons of the Vth Republic. PCF refers to Marchais much less often than it does Mitterrand.

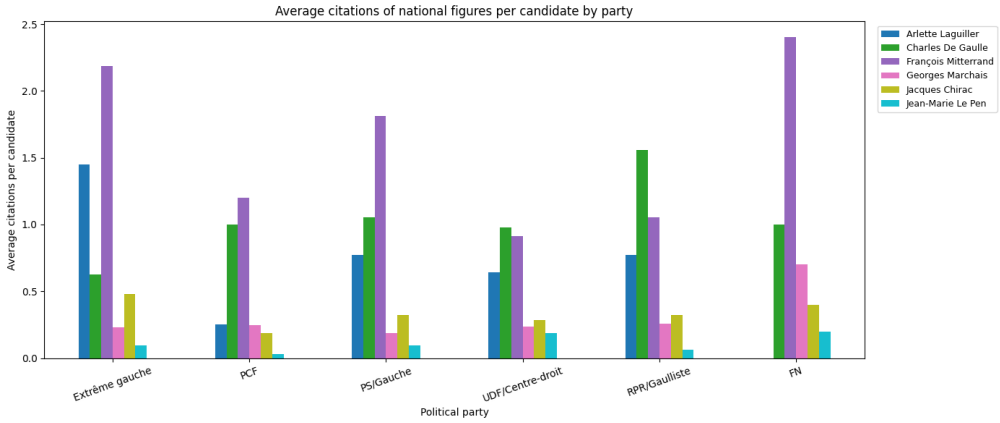

Figure 3: Average number of citations of national figures per candidates by political party ordered left to right on the ideological spectrum. Mitterrand dominates across all parties; Marchais is underrepresented even within its own party the PCF.

5.3 Alignment vs. distancing

Alignment scores by party and figure, sourced from Table2 & Figure4. Several patterns are historically coherent: the Extreme gauche cites De Gaulle negatively (-1,00), Marchais finds himself condemned across the board including by PCF itself (-0.52) and RPR distance Le Pen (- 0,33). Well-documented strategic ambiguity characterize Chirac in 1981, and explains the RPR placement on Mitterrand (0.63). The PCF's aberrantly positive score on Le Pen (0.60)(For methodological reasons, see Section 6.)

Table 2: Lexical alignment proxy by party and figure (+1 = positive lexical context, -1 = negative lexical context).

Party	Mitterrand	De Gaulle	Marchais	Chirac	Le Pen
Extreme gauche	0.45	-1.00	-0.46	0.44	0.00
PCF	0.69	0.29	-0.52	0.33	0.60
PS/Gauche	0.51	0.20	-0.50	0.56	0.00
UDF/Centre-droit	0.45	0.33	-1.00	0.57	0.00
RPR/Gaulliste	0.63	0.27	-1.00	0.43	-0.33
FN	0.60	0.00	0.00	-1.00	0.00

Scores based on few marker-bearing contexts should be interpreted with caution, especially when mentions are administrative rather than rhetorical.

5.4 Geographic scope

So far, national references consist 65% of all classified geographic mentions, compared with 8% local and 25% international references. The evolution by year can be seen in Figure 5. One remarkable pattern is the increase of international mentions after 1988, which we explain with increased salience of the Maastricht debate. Europe leads international references (3,295 mentions) in front of Algeria (425), Germany coming third with 320.

Indeed, geographic scope also varies by party with an extremely high share of international references (37%) for FN, presumably reflecting anti-European and anti-immigrant rhetoric, and Divers centre candidates overwhelmingly framing issues nationally. On regional basis, the Haut-Rhin has a local reference share of 22% and thus is not only much higher than other departments in terms of identity at a regional level.

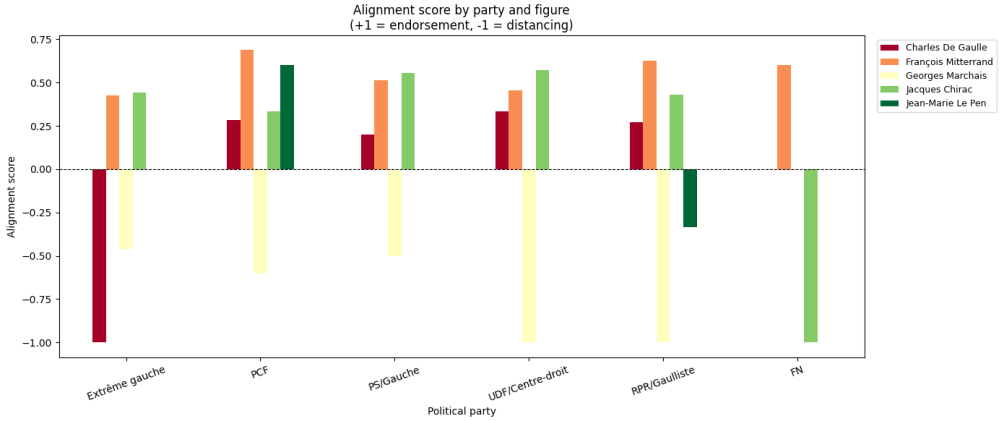

Figure 4: Alignment scores by party and figure. Positive values indicate endorsement, negative values distancing. Marchais is consistently distanced; Mitterrand receives positive scores from most parties.

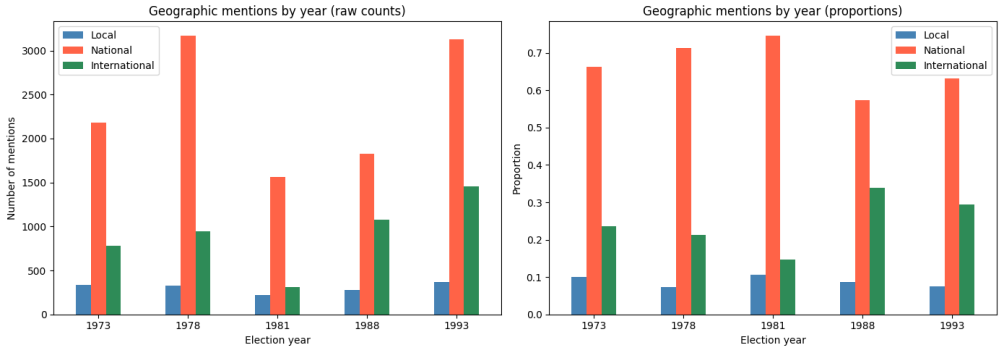

Figure 5: Geographic scope of the discourses by election year (represented in proportions). National references dominate and the international references rise from 1988, consistent with the Maastricht ratification debate of 1992.

5.5 Are the two dimensions linked?

The Pearson correlation of national share and citation count are as follow : Mitterrand at 0.028, De Gaulle at 0.023 and Le Pen at -0.082. This suggests that, in this corpus, the two rhetorical dimensions are weakly associated and can be studied separately.

6 Discussion

And that is exactly what we find: the national references dominate the classified geographic mentions in this corpus throughout the period. Irrespective of their political allegiance, candidates reference national figures and present local politics through national-level stakes, as they do with little geographic variation apart from regions beholden to solid identities like Alsace. The centrality of Mitterrand within the political world is easily visible, by both his raw citation numbers and by how non-partisan those citations are.

Certainly the context window analysis has its limits. Lexicon-level sentiment is unable to use context of negation, irony or administrative citation. This is evident in the PCF's anomalous positive score on Le Pen (0.60): most mentions of Le Pen in PCF manifestos are merely administrative listings against other candidates — our lexicon cannot differentiate this from a rhetorical engagement. To determine this reliably, one would need a fine tuned sequence classifier or similar sentiment model trained on annotated electoral text.

The results of the geographic analysis are limited by the design of the corpus. A mere five departments (3 peripheral [Alsace, Pyrenees]). Therefore, the results should not be interpreted as representative of all French legislative campaigns. To make more definitive statements about the balance local-national in French political discourse generally, a corpus including all metropolitan departments would be necessary.

7 Conclusion

In this paper Named Entity Recognition was applied to the Archelec corpus so as to look into two dimensions, over a 20-year history of French politics. Mitterrand's across-the-board dominance on the sponsorship dimension is a reflection of his singular status at the heart of French politics from 1973 to 1993. The alignment analysis produces patterns that are partly consistent with historical expectations, while also highlighting a key limitation of lexicon-based sentiment applied to political citation contexts. National framing dominates over local throughout the spatial dimension whilst international references rise from 1988 in sync with the Maastricht debate. The empirical independence of the two dimensions suggests that invoking national figures and projecting nationally in geographic terms are different rhetorical strategies which require separate study. These findings should nevertheless be read as exploratory, given OCR noise, limited geographic coverage and the absence of domain-specific NER annotation.

References

Gavin Abercrombie and Riza Batista-Navarro. Sentiment and position-taking in political discourse: A survey of approaches. In *Proceedings of the Workshop on Computational Approaches to Subjectivity, Sentiment and Social Media Analysis*, 2020.

Ian Budge and Hans-Dieter Klingemann. *Mapping Policy Preferences*. Oxford University Press, 2001.

Maud Ehrmann, Ahmed Hamdi, Elvys Linhares Pontes, Matteo Romanello, and Antoine Doucet. Named entity recognition and classification in historical documents: A survey. *ACM Computing Surveys*, 56(2), 2023.

Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. spacy: Industrial-strength natural language processing in python, 2020.

Kevin M. Quinn, Burt L. Monroe, Michael Colaresi, Michael H. Crespin, and Dragomir R. Radev. How to analyze political attention with minimal assumptions and costs. *American Journal of Political Science*, 54(1):209–228, 2010.

Ludovic Rheault and Christopher Cochrane. Word embeddings for the analysis of ideological placement in parliamentary corpora. *Political Analysis*, 28(1):112–133, 2020.

G. Vidal et al. Adminset: A french administrative ocr corpus for ner. *HAL*, 2024. hal-04855066.

Urchade Zaratiana, Nadi Tomeh, Pierre Holat, and Thierry Charnois. Gliner: Generalist model for named entity recognition using bidirectional transformer. *arXiv*, 2023. arXiv:2311.08526.